

4

ARM Instruction Set

This chapter describes the ARM instruction set.

| | | |
|------|---|------|
| 4.1 | Instruction Set Summary | 4-2 |
| 4.2 | The Condition Field | 4-5 |
| 4.4 | Branch and Branch with Link (B, BL) | 4-8 |
| 4.5 | Data Processing | 4-10 |
| 4.7 | Multiply and Multiply-Accumulate (MUL, MLA) | 4-23 |
| 4.9 | Single Data Transfer (LDR, STR) | 4-28 |
| 4.11 | Block Data Transfer (LDM, STM) | 4-40 |
| 4.13 | Software Interrupt (SWI) | 4-49 |
| 4.17 | Undefined Instruction | 4-60 |

ARM Instruction Set - Summary

4.1 Instruction Set Summary

4.1.1 Format summary

The ARM instruction set formats are shown below.

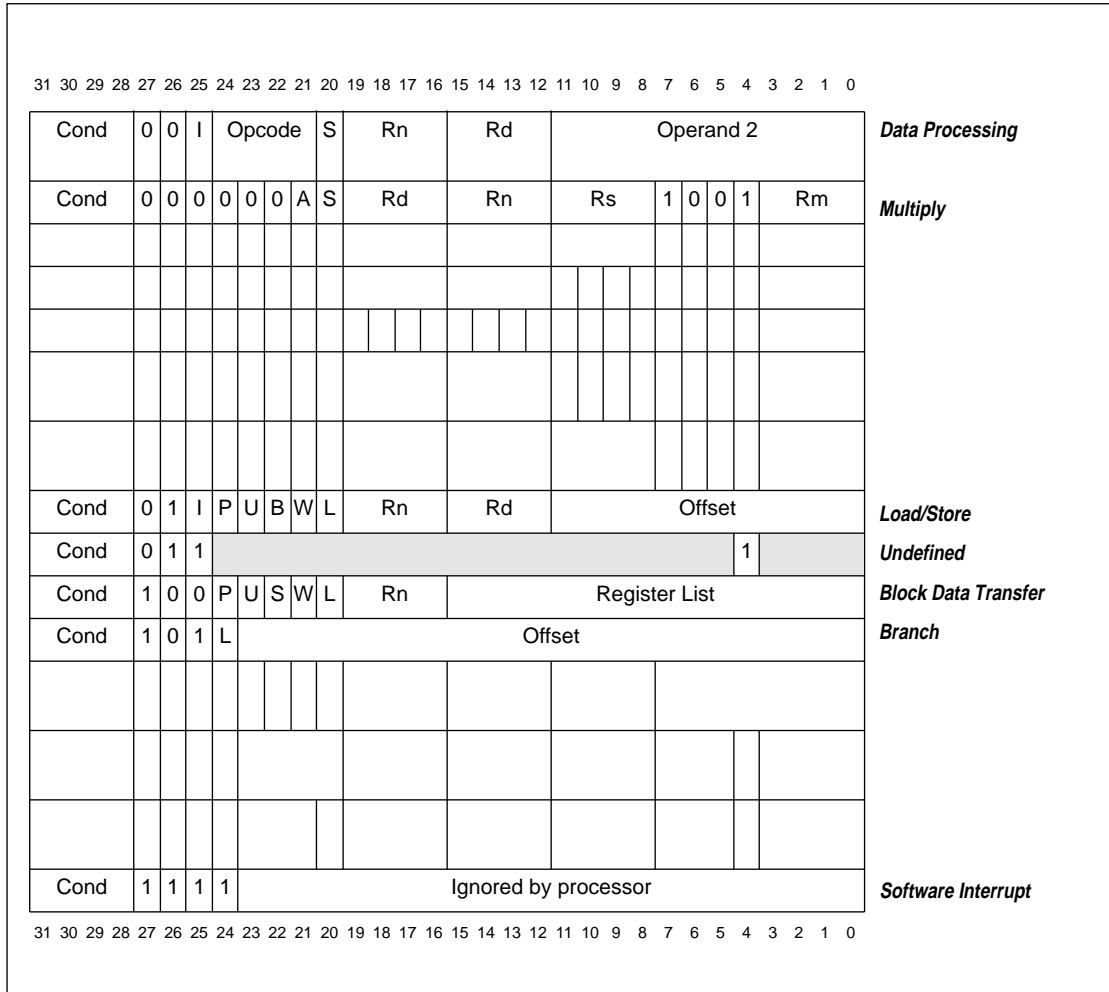


Figure 4-1: ARM instruction set formats

ARM Instruction Set - Summary

4.1.2 Instruction summary

| Mnemonic | Instruction | Action | See Section: |
|----------|---------------------------|---|--------------|
| ADC | Add with carry | $Rd := Rn + Op2 + \text{Carry}$ | 4.5 |
| ADD | Add | $Rd := Rn + Op2$ | 4.5 |
| AND | AND | $Rd := Rn \text{ AND } Op2$ | 4.5 |
| B | Branch | $R15 := \text{address}$ | 4.4 |
| BIC | Bit Clear | $Rd := Rn \text{ AND NOT } Op2$ | 4.5 |
| BL | Branch with Link | $R14 := R15, R15 := \text{address}$ | 4.4 |
| | | | |
| CMN | Compare Negative | $\text{CPSR flags} := Rn + Op2$ | 4.5 |
| CMP | Compare | $\text{CPSR flags} := Rn - Op2$ | 4.5 |
| EOR | Exclusive OR | $Rd := (Rn \text{ AND NOT } Op2) \text{ OR } (Op2 \text{ AND NOT } Rn)$ | 4.5 |
| | | | |
| LDM | Load multiple registers | Stack manipulation (Pop) | 4.11 |
| LDR | Load register from memory | $Rd := (\text{address})$ | 4.9, 4.10 |
| | | | |
| MLA | Multiply Accumulate | $Rd := (Rm * Rs) + Rn$ | 4.7, 4.8 |
| MOV | Move register or constant | $Rd := Op2$ | 4.5 |
| | | | |
| | | | |
| MUL | Multiply | $Rd := Rm * Rs$ | 4.7, 4.8 |
| MVN | Move negative register | $Rd := 0xFFFFFFFF \text{ EOR } Op2$ | 4.5 |
| ORR | OR | $Rd := Rn \text{ OR } Op2$ | 4.5 |

Table 4-1: The ARM Instruction set

ARM Instruction Set - Summary

| Mnemonic | Instruction | Action | See Section: |
|----------|-----------------------------|--|--------------|
| RSB | Reverse Subtract | $Rd := Op2 - Rn$ | 4.5 |
| RSC | Reverse Subtract with Carry | $Rd := Op2 - Rn - 1 + \text{Carry}$ | 4.5 |
| SBC | Subtract with Carry | $Rd := Rn - Op2 - 1 + \text{Carry}$ | 4.5 |
| | | | |
| STM | Store Multiple | Stack manipulation (Push) | 4.11 |
| STR | Store register to memory | $\langle \text{address} \rangle := Rd$ | 4.9, 4.10 |
| SUB | Subtract | $Rd := Rn - Op2$ | 4.5 |
| SWI | Software Interrupt | OS call | 4.13 |
| | | | |
| TEQ | Test bitwise equality | $\text{CPSR flags} := Rn \text{ EOR } Op2$ | 4.5 |
| TST | Test bits | $\text{CPSR flags} := Rn \text{ AND } Op2$ | 4.5 |

Table 4-1: The ARM Instruction set (Continued)

4.2 The Condition Field

In ARM state, all instructions are conditionally executed according to the state of the CPSR condition codes and the instruction's condition field. This field (bits 31:28) determines the circumstances under which an instruction is to be executed. If the state of the C, N, Z and V flags fulfils the conditions encoded by the field, the instruction is executed, otherwise it is ignored.

There are sixteen possible conditions, each represented by a two-character suffix that can be appended to the instruction's mnemonic. For example, a Branch (B in assembly language) becomes BEQ for "Branch if Equal", which means the Branch will only be taken if the Z flag is set.

In practice, fifteen different conditions may be used: these are listed in [Table 4-2: Condition code summary](#). The sixteenth (1111) is reserved, and must not be used.

In the absence of a suffix, the condition field of most instructions is set to "Always" (suffix AL). This means the instruction will always be executed regardless of the CPSR condition codes.

| Code | Suffix | Flags | Meaning |
|------|--------|-----------------------------|-------------------------|
| 0000 | EQ | Z set | equal |
| 0001 | NE | Z clear | not equal |
| 0010 | CS | C set | unsigned higher or same |
| 0011 | CC | C clear | unsigned lower |
| 0100 | MI | N set | negative |
| 0101 | PL | N clear | positive or zero |
| 0110 | VS | V set | overflow |
| 0111 | VC | V clear | no overflow |
| 1000 | HI | C set and Z clear | unsigned higher |
| 1001 | LS | C clear or Z set | unsigned lower or same |
| 1010 | GE | N equals V | greater or equal |
| 1011 | LT | N not equal to V | less than |
| 1100 | GT | Z clear AND (N equals V) | greater than |
| 1101 | LE | Z set OR (N not equal to V) | less than or equal |
| 1110 | AL | (ignored) | always |

Table 4-2: Condition code summary

ARM Instruction Set - B, BL

4.4 Branch and Branch with Link (B, BL)

The instruction is only executed if the condition is true. The various conditions are defined in [Table 4-2: Condition code summary](#) on page 4-5. The instruction encoding is shown in [Figure 4-3: Branch instructions](#), below.




Figure 4-3: Branch instructions

Branch instructions contain a signed 2's complement 24 bit offset. This is shifted left two bits, sign extended to 32 bits, and added to the PC. The instruction can therefore specify a branch of +/- 32Mbytes. The branch offset must take account of the prefetch operation, which causes the PC to be 2 words (8 bytes) ahead of the current instruction.

Branches beyond +/- 32Mbytes must use an offset or absolute destination which has been previously loaded into a register. In this case the PC should be manually saved in R14 if a Branch with Link type operation is required.

4.4.1 The link bit

Branch with Link (BL) writes the old PC into the link register (R14) of the current bank. The PC value written into R14 is adjusted to allow for the prefetch, and contains the address of the instruction following the branch and link instruction. Note that the CPSR is not saved with the PC and R14[1:0] are always cleared.

To return from a routine called by Branch with Link use `MOV PC,R14` if the link register is still valid or `LDM Rn!,{..PC}` if the link register has been saved onto a stack pointed to by Rn.

4.4.2 Instruction cycle times

Branch and Branch with Link instructions take $2S + 1N$ incremental cycles, where S and N are as defined in [6.2 Cycle Types](#) on page 6-2.

4.4.3 Assembler syntax

Items in {} are optional. Items in <> must be present.

B{L}{cond} <expression>

{L} is used to request the Branch with Link form of the instruction. If absent, R14 will not be affected by the instruction.

{cond} is a two-character mnemonic as shown in [Table 4-2: Condition code summary](#) on page 4-5. If absent then AL (ALways) will be used.

<expression> is the destination. The assembler calculates the offset.

4.4.4 Examples

```
here  BAL  here    ; assembles to 0xEAFFFFFEE (note effect of
                  ; PC offset).
      B    there   ; Always condition used as default.
      CMP  R1,#0    ; Compare R1 with zero and branch to fred
                  ; if R1 was zero, otherwise continue
      BEQ  fred     ; continue to next instruction.

      BL   sub+ROM   ; Call subroutine at computed address.
      ADDS R1,#1     ; Add 1 to register 1, setting CPSR flags
                  ; on the result then call subroutine if
      BLCC sub       ; the C flag is clear, which will be the
                  ; case unless R1 held 0xFFFFFFFF.
```

ARM Instruction Set - Data processing

4.5 Data Processing

The data processing instruction is only executed if the condition is true. The conditions are defined in **Table 4-2: Condition code summary** on page 4-5.

The instruction encoding is shown in **Figure 4-4: Data processing instructions** below.




Figure 4-4: Data processing instructions

The instruction produces a result by performing a specified arithmetic or logical operation on one or two operands. The first operand is always a register (Rn).

ARM Instruction Set - Data processing

The second operand may be a shifted register (Rm) or a rotated 8 bit immediate value (Imm) according to the value of the I bit in the instruction. The condition codes in the CPSR may be preserved or updated as a result of this instruction, according to the value of the S bit in the instruction.

Certain operations (TST, TEQ, CMP, CMN) do not write the result to Rd. They are used only to perform tests and to set the condition codes on the result and always have the S bit set. The instructions and their effects are listed in **Table 4-3: ARM Data processing instructions** on page 4-11.

4.5.1 CPSR flags

The data processing operations may be classified as logical or arithmetic. The logical operations (AND, EOR, TST, TEQ, ORR, MOV, BIC, MVN) perform the logical action on all corresponding bits of the operand or operands to produce the result. If the S bit is set (and Rd is not R15, see below) the V flag in the CPSR will be unaffected, the C flag will be set to the carry out from the barrel shifter (or preserved when the shift operation is LSL #0), the Z flag will be set if and only if the result is all zeros, and the N flag will be set to the logical value of bit 31 of the result.

| Assembler Mnemonic | OpCode | Action |
|--------------------|--------|---------------------------------------|
| AND | 0000 | operand1 AND operand2 |
| EOR | 0001 | operand1 EOR operand2 |
| SUB | 0010 | operand1 - operand2 |
| RSB | 0011 | operand2 - operand1 |
| ADD | 0100 | operand1 + operand2 |
| ADC | 0101 | operand1 + operand2 + carry |
| SBC | 0110 | operand1 - operand2 + carry - 1 |
| RSC | 0111 | operand2 - operand1 + carry - 1 |
| TST | 1000 | as AND, but result is not written |
| TEQ | 1001 | as EOR, but result is not written |
| CMP | 1010 | as SUB, but result is not written |
| CMN | 1011 | as ADD, but result is not written |
| ORR | 1100 | operand1 OR operand2 |
| MOV | 1101 | operand2 (operand1 is ignored) |
| BIC | 1110 | operand1 AND NOT operand2 (Bit clear) |
| MVN | 1111 | NOT operand2 (operand1 is ignored) |

Table 4-3: ARM Data processing instructions

ARM Instruction Set - Shifts

The arithmetic operations (SUB, RSB, ADD, ADC, SBC, RSC, CMP, CMN) treat each operand as a 32 bit integer (either unsigned or 2's complement signed, the two are equivalent). If the S bit is set (and Rd is not R15) the V flag in the CPSR will be set if an overflow occurs into bit 31 of the result; this may be ignored if the operands were considered unsigned, but warns of a possible error if the operands were 2's complement signed. The C flag will be set to the carry out of bit 31 of the ALU, the Z flag will be set if and only if the result was zero, and the N flag will be set to the value of bit 31 of the result (indicating a negative result if the operands are considered to be 2's complement signed).

4.5.2 Shifts

When the second operand is specified to be a shifted register, the operation of the barrel shifter is controlled by the Shift field in the instruction. This field indicates the type of shift to be performed (logical left or right, arithmetic right or rotate right). The amount by which the register should be shifted may be contained in an immediate field in the instruction, or in the bottom byte of another register (other than R15). The encoding for the different shift types is shown in **Figure 4-5: ARM shift operations**.




Figure 4-5: ARM shift operations

Instruction specified shift amount

When the shift amount is specified in the instruction, it is contained in a 5 bit field which may take any value from 0 to 31. A logical shift left (LSL) takes the contents of Rm and moves each bit by the specified amount to a more significant position. The least significant bits of the result are filled with zeros, and the high bits of Rm which do not map into the result are discarded, except that the least significant discarded bit becomes the shifter carry output which may be latched into the C bit of the CPSR when the ALU operation is in the logical class (see above). For example, the effect of LSL #5 is shown in **Figure 4-6: Logical shift left**.

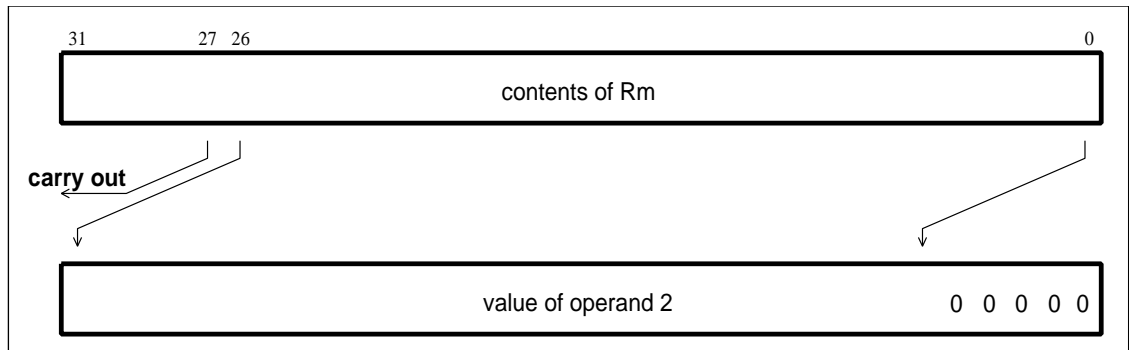


Figure 4-6: Logical shift left

Note LSL #0 is a special case, where the shifter carry out is the old value of the CPSR C flag. The contents of Rm are used directly as the second operand.

A logical shift right (LSR) is similar, but the contents of Rm are moved to less significant positions in the result. LSR #5 has the effect shown in **Figure 4-7: Logical shift right**.




Figure 4-7: Logical shift right

The form of the shift field which might be expected to correspond to LSR #0 is used to encode LSR #32, which has a zero result with bit 31 of Rm as the carry output. Logical shift right zero is redundant as it is the same as logical shift left zero, so the assembler will convert LSR #0 (and ASR #0 and ROR #0) into LSL #0, and allow LSR #32 to be specified.

An arithmetic shift right (ASR) is similar to logical shift right, except that the high bits are filled with bit 31 of Rm instead of zeros. This preserves the sign in 2's complement notation. For example, ASR #5 is shown in **Figure 4-8: Arithmetic shift right**.

ARM Instruction Set - Shifts

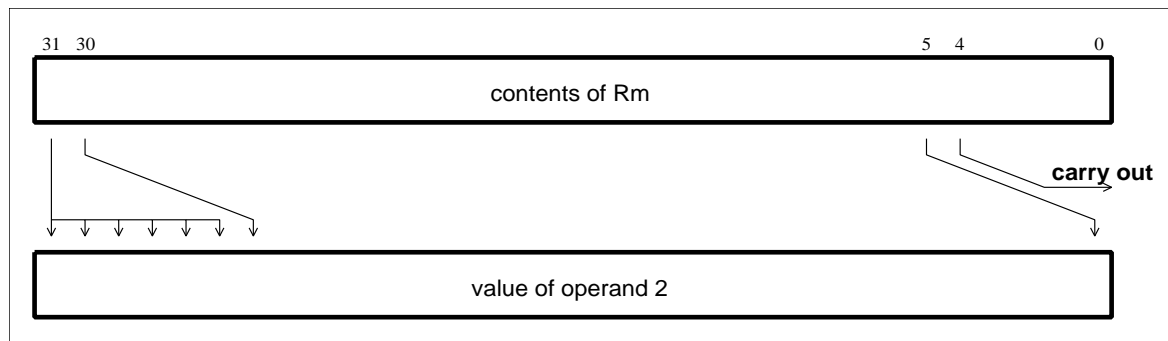


Figure 4-8: Arithmetic shift right

The form of the shift field which might be expected to give ASR #0 is used to encode ASR #32. Bit 31 of Rm is again used as the carry output, and each bit of operand 2 is also equal to bit 31 of Rm. The result is therefore all ones or all zeros, according to the value of bit 31 of Rm.

Rotate right (ROR) operations reuse the bits which “overshoot” in a logical shift right operation by reintroducing them at the high end of the result, in place of the zeros used to fill the high end in logical right operations. For example, ROR #5 is shown in [Figure 4-9: Rotate right](#) on page 4-14.

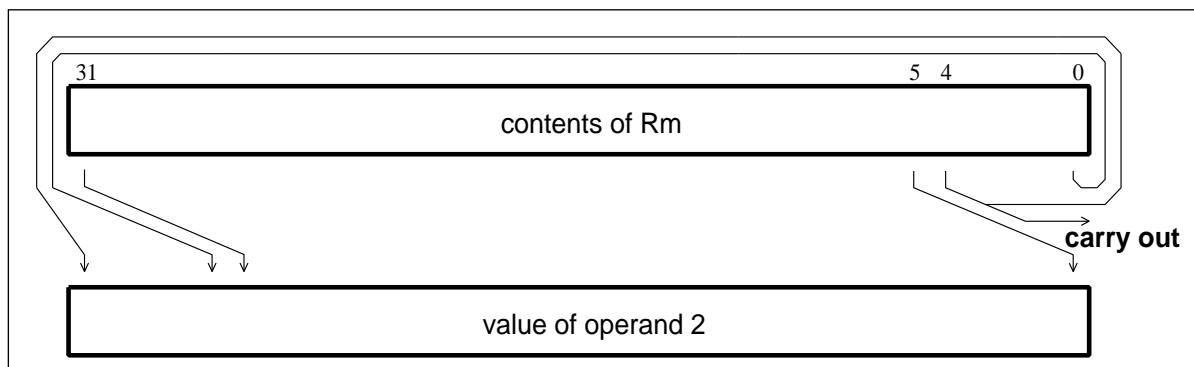


Figure 4-9: Rotate right

The form of the shift field which might be expected to give ROR #0 is used to encode a special function of the barrel shifter, rotate right extended (RRX). This is a rotate right by one bit position of the 33 bit quantity formed by appending the CPSR C flag to the most significant end of the contents of Rm as shown in [Figure 4-10: Rotate right extended](#).




Figure 4-10: Rotate right extended

Register specified shift amount

Only the least significant byte of the contents of Rs is used to determine the shift amount. Rs can be any general register other than R15.

If this byte is zero, the unchanged contents of Rm will be used as the second operand, and the old value of the CPSR C flag will be passed on as the shifter carry output.

If the byte has a value between 1 and 31, the shifted result will exactly match that of an instruction specified shift with the same value and shift operation.

If the value in the byte is 32 or more, the result will be a logical extension of the shift described above:

- 1 LSL by 32 has result zero, carry out equal to bit 0 of Rm.
- 2 LSL by more than 32 has result zero, carry out zero.
- 3 LSR by 32 has result zero, carry out equal to bit 31 of Rm.
- 4 LSR by more than 32 has result zero, carry out zero.
- 5 ASR by 32 or more has result filled with and carry out equal to bit 31 of Rm.
- 6 ROR by 32 has result equal to Rm, carry out equal to bit 31 of Rm.
- 7 ROR by n where n is greater than 32 will give the same result and carry out as ROR by n-32; therefore repeatedly subtract 32 from n until the amount is in the range 1 to 32 and see above.

Note The zero in bit 7 of an instruction with a register controlled shift is compulsory; a one in this bit will cause the instruction to be a multiply or undefined instruction.

4.5.3 Immediate operand rotates

The immediate operand rotate field is a 4 bit unsigned integer which specifies a shift operation on the 8 bit immediate value. This value is zero extended to 32 bits, and then subject to a rotate right by twice the value in the rotate field. This enables many common constants to be generated, for example all powers of 2.

ARM Instruction Set - TEQ, TST, CMP & CMN

4.5.4 Writing to R15

When Rd is a register other than R15, the condition code flags in the CPSR may be updated from the ALU flags as described above.

When Rd is R15 and the S flag in the instruction is not set the result of the operation is placed in R15 and the CPSR is unaffected.

When Rd is R15 and the S flag is set the result of the operation is placed in R15 and the SPSR corresponding to the current mode is moved to the CPSR. This allows state changes which atomically restore both PC and CPSR. This form of instruction should not be used in User mode.

4.5.5 Using R15 as an operand

If R15 (the PC) is used as an operand in a data processing instruction the register is used directly.

The PC value will be the address of the instruction, plus 8 or 12 bytes due to instruction prefetching. If the shift amount is specified in the instruction, the PC will be 8 bytes ahead. If a register is used to specify the shift amount the PC will be 12 bytes ahead.

4.5.6 TEQ, TST, CMP and CMN opcodes

Note *TEQ, TST, CMP and CMN do not write the result of their operation but do set flags in the CPSR. An assembler should always set the S flag for these instructions even if this is not specified in the mnemonic.*

The TEQP form of the TEQ instruction used in earlier ARM processors must not be used: the PSR transfer operations should be used instead.

The action of TEQP in the ARM7TDMI is to move SPSR_<mode> to the CPSR if the processor is in a privileged mode and to do nothing if in User mode.

4.5.7 Instruction cycle times

Data Processing instructions vary in the number of incremental cycles taken as follows:

| Processing Type | Cycles |
|--|--------------|
| Normal Data Processing | 1S |
| Data Processing with register specified shift | 1S + 1I |
| Data Processing with PC written | 2S + 1N |
| Data Processing with register specified shift and PC written | 2S + 1N + 1I |

Table 4-4: Incremental cycle times

S, N and I are as defined in [6.2 Cycle Types](#) on page 6-2.

ARM Instruction Set - TEQ, TST, CMP & CMN

4.5.8 Assembler syntax

- 1 MOV,MVN (single operand instructions.)
<opcode>{cond}{S} Rd,<Op2>
- 2 CMP,CMN,TEQ,TST (instructions which do not produce a result.)
<opcode>{cond} Rn,<Op2>
- 3 AND,EOR,SUB,RSB,ADD,ADC,SBC,RSC,ORR,BIC
<opcode>{cond}{S} Rd,Rn,<Op2>

where:

| | |
|---------------|--|
| <Op2> | is Rm{,<shift>} or,<#expression> |
| {cond} | is a two-character condition mnemonic. See Table 4-2: Condition code summary on page 4-5. |
| {S} | set condition codes if S present (implied for CMP, CMN, TEQ, TST). |
| Rd, Rn and Rm | are expressions evaluating to a register number. |
| <#expression> | if this is used, the assembler will attempt to generate a shifted immediate 8-bit field to match the expression. If this is impossible, it will give an error. |
| <shift> | is <shiftname> <register> or <shiftname> #expression, or RRX (rotate right one bit with extend). |
| <shiftname>s | are: ASL, LSL, LSR, ASR, ROR. (ASL is a synonym for LSL, they assemble to the same code.) |

4.5.9 Examples

```
ADDEQ R2,R4,R5      ; If the Z flag is set make R2:=R4+R5
TEQS  R4,#3          ; test R4 for equality with 3.
                    ; (The S is in fact redundant as the
                    ; assembler inserts it automatically.)
SUB    R4,R5,R7,LSR R2; Logical right shift R7 by the number in
                    ; the bottom byte of R2, subtract result
                    ; from R5, and put the answer into R4.
MOV    PC,R14        ; Return from subroutine.
MOVS   PC,R14        ; Return from exception and restore CPSR
                    ; from SPSR_mode.
```



4.7 Multiply and Multiply-Accumulate (MUL, MLA)

The instruction is only executed if the condition is true. The various conditions are defined in [Table 4-2: Condition code summary](#) on page 4-5. The instruction encoding is shown in [Figure 4-12: Multiply instructions](#).

The multiply and multiply-accumulate instructions use an 8 bit Booth's algorithm to perform integer multiplication.




Figure 4-12: Multiply instructions

The multiply form of the instruction gives $Rd := Rm * Rs$. Rn is ignored, and should be set to zero for compatibility with possible future upgrades to the instruction set.

The multiply-accumulate form gives $Rd := Rm * Rs + Rn$, which can save an explicit ADD instruction in some circumstances.

Both forms of the instruction work on operands which may be considered as signed (2's complement) or unsigned integers.

The results of a signed multiply and of an unsigned multiply of 32 bit operands differ only in the upper 32 bits - the low 32 bits of the signed and unsigned results are identical. As these instructions only produce the low 32 bits of a multiply, they can be used for both signed and unsigned multiplies.

For example consider the multiplication of the operands:

| Operand A | Operand B | Result |
|-------------|-----------|--------------|
| 0xFFFFFFFF6 | 0x0000001 | 0xFFFFFFFF38 |

If the operands are interpreted as signed

Operand A has the value -10, operand B has the value 20, and the result is -200 which is correctly represented as 0xFFFFFFFF38

If the operands are interpreted as unsigned

Operand A has the value 4294967286, operand B has the value 20 and the result is 85899345720, which is represented as 0x13FFFFFF38, so the least significant 32 bits are 0xFFFFFFFF38.

ARM Instruction Set - MUL, MLA

4.7.1 Operand restrictions

The destination register Rd must not be the same as the operand register Rm. R15 must not be used as an operand or as the destination register.

All other register combinations will give correct results, and Rd, Rn and Rs may use the same register when required.

4.7.2 CPSR flags

Setting the CPSR flags is optional, and is controlled by the S bit in the instruction. The N (Negative) and Z (Zero) flags are set correctly on the result (N is made equal to bit 31 of the result, and Z is set if and only if the result is zero). The C (Carry) flag is set to a meaningless value and the V (oVerflow) flag is unaffected.

4.7.3 Instruction cycle times

MUL takes $1S + mI$ and MLA $1S + (m+1)I$ cycles to execute, where S and I are as defined in [C6.2 Cycle Types](#) on page 6-2.

| | |
|---|--|
| m | is the number of 8 bit multiplier array cycles required to complete the multiply, which is controlled by the value of the multiplier operand specified by Rs. Its possible values are as follows |
| 1 | if bits [32:8] of the multiplier operand are all zero or all one. |
| 2 | if bits [32:16] of the multiplier operand are all zero or all one. |
| 3 | if bits [32:24] of the multiplier operand are all zero or all one. |
| 4 | in all other cases. |

4.7.4 Assembler syntax

| | |
|-------------------|--|
| MUL{cond}{S} | Rd,Rm,Rs |
| MLA{cond}{S} | Rd,Rm,Rs,Rn |
| {cond} | two-character condition mnemonic. See Table 4-2: Condition code summary on page 4-5. |
| {S} | set condition codes if S present |
| Rd, Rm, Rs and Rn | are expressions evaluating to a register number other than R15. |

4.7.5 Examples

```
MUL      R1,R2,R3      ; R1:=R2*R3
MLAEQS   R1,R2,R3,R4   ; Conditionally R1:=R2*R3+R4,
                        ; setting condition codes.
```

ARM Instruction Set - LDR, STR

4.9 Single Data Transfer (LDR, STR)

The instruction is only executed if the condition is true. The various conditions are defined in [Table 4-2: Condition code summary](#) on page 4-5. The instruction encoding is shown in [Figure 4-14: Single data transfer instructions](#) on page 4-28.

The single data transfer instructions are used to load or store single bytes or words of data. The memory address used in the transfer is calculated by adding an offset to or subtracting an offset from a base register.

The result of this calculation may be written back into the base register if auto-indexing is required.

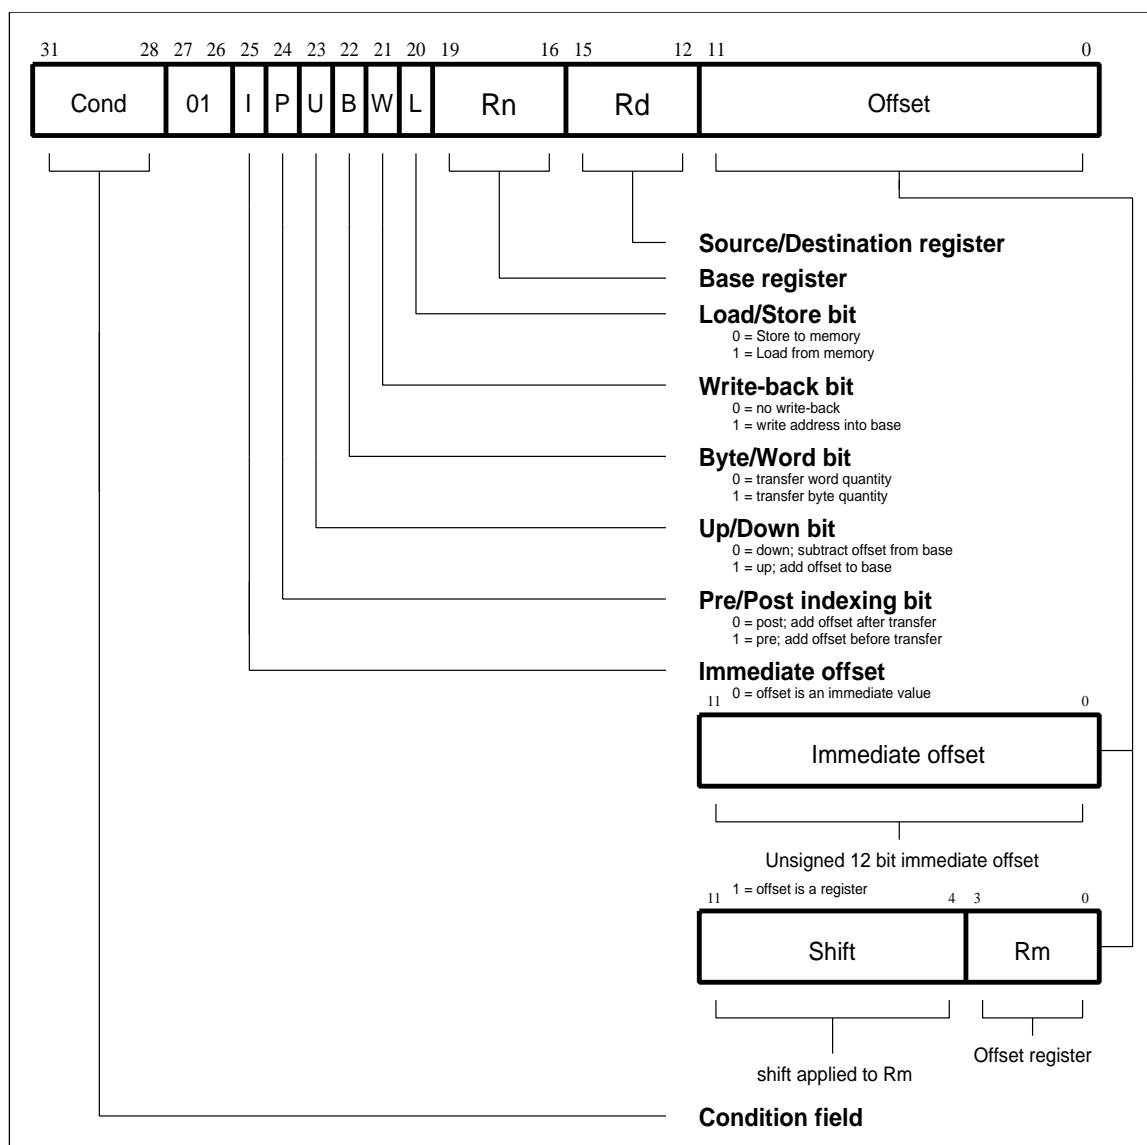


Figure 4-14: Single data transfer instructions

4.9.1 Offsets and auto-indexing

The offset from the base may be either a 12 bit unsigned binary immediate value in the instruction, or a second register (possibly shifted in some way). The offset may be added to (U=1) or subtracted from (U=0) the base register Rn. The offset modification may be performed either before (pre-indexed, P=1) or after (post-indexed, P=0) the base is used as the transfer address.

The W bit gives optional auto increment and decrement addressing modes. The modified base value may be written back into the base (W=1), or the old base value may be kept (W=0). In the case of post-indexed addressing, the write back bit is redundant and is always set to zero, since the old base value can be retained by setting the offset to zero. Therefore post-indexed data transfers always write back the modified base. The only use of the W bit in a post-indexed data transfer is in privileged mode code, where setting the W bit forces non-privileged mode for the transfer, allowing the operating system to generate a user address in a system where the memory management hardware makes suitable use of this hardware.

4.9.2 Shifted register offset

The 8 shift control bits are described in the data processing instructions section. However, the register specified shift amounts are not available in this instruction class. See [4.5.2 Shifts](#) on page 4-12.

4.9.3 Bytes and words

This instruction class may be used to transfer a byte (B=1) or a word (B=0) between an ARM7TDMI register and memory.

The action of LDR(B) and STR(B) instructions is influenced by the **BIGEND** control signal. The two possible configurations are described below.

Little endian configuration

A byte load (LDRB) expects the data on data bus inputs 7 through 0 if the supplied address is on a word boundary, on data bus inputs 15 through 8 if it is a word address plus one byte, and so on. The selected byte is placed in the bottom 8 bits of the destination register, and the remaining bits of the register are filled with zeros. Please see [Figure 3-2: Little endian addresses of bytes within words](#) on page 3-3.

A byte store (STRB) repeats the bottom 8 bits of the source register four times across data bus outputs 31 through 0. The external memory system should activate the appropriate byte subsystem to store the data.

A word load (LDR) will normally use a word aligned address. However, an address offset from a word boundary will cause the data to be rotated into the register so that the addressed byte occupies bits 0 to 7. This means that half-words accessed at offsets 0 and 2 from the word boundary will be correctly loaded into bits 0 through 15 of the register. Two shift operations are then required to clear or to sign extend the upper 16 bits. This is illustrated in [Figure 4-15: Little endian offset addressing](#) on page 4-30.

ARM Instruction Set - LDR, STR




Figure 4-15: Little endian offset addressing

A word store (STR) should generate a word aligned address. The word presented to the data bus is not affected if the address is not word aligned. That is, bit 31 of the register being stored always appears on data bus output 31.

Big endian configuration

A byte load (LDRB) expects the data on data bus inputs 31 through 24 if the supplied address is on a word boundary, on data bus inputs 23 through 16 if it is a word address plus one byte, and so on. The selected byte is placed in the bottom 8 bits of the destination register and the remaining bits of the register are filled with zeros. Please see [Figure 3-1: Big endian addresses of bytes within words](#) on page 3-3.

A byte store (STRB) repeats the bottom 8 bits of the source register four times across data bus outputs 31 through 0. The external memory system should activate the appropriate byte subsystem to store the data.

A word load (LDR) should generate a word aligned address. An address offset of 0 or 2 from a word boundary will cause the data to be rotated into the register so that the addressed byte occupies bits 31 through 24. This means that half-words accessed at these offsets will be correctly loaded into bits 16 through 31 of the register. A shift operation is then required to move (and optionally sign extend) the data into the bottom 16 bits. An address offset of 1 or 3 from a word boundary will cause the data to be rotated into the register so that the addressed byte occupies bits 15 through 8.

A word store (STR) should generate a word aligned address. The word presented to the data bus is not affected if the address is not word aligned. That is, bit 31 of the register being stored always appears on data bus output 31.

4.9.4 Use of R15

Write-back must not be specified if R15 is specified as the base register (Rn). When using R15 as the base register you must remember it contains an address 8 bytes on from the address of the current instruction.

R15 must not be specified as the register offset (Rm).

When R15 is the source register (Rd) of a register store (STR) instruction, the stored value will be address of the instruction plus 12.

4.9.5 Restriction on the use of base register

When configured for late aborts, the following example code is difficult to unwind as the base register, Rn, gets updated before the abort handler starts. Sometimes it may be impossible to calculate the initial value.

After an abort, the following example code is difficult to unwind as the base register, Rn, gets updated before the abort handler starts. Sometimes it may be impossible to calculate the initial value.

Example:

```
LDR    R0, [R1], R1
```

Therefore a post-indexed LDR or STR where Rm is the same register as Rn should not be used.

4.9.6 Data aborts

A transfer to or from a legal address may cause problems for a memory management system. For instance, in a system which uses virtual memory the required data may be absent from main memory. The memory manager can signal a problem by taking the processor **ABORT** input HIGH whereupon the Data Abort trap will be taken. It is up to the system software to resolve the cause of the problem, then the instruction can be restarted and the original program continued.

4.9.7 Instruction cycle times

Normal LDR instructions take $1S + 1N + 1I$ and LDR PC take $2S + 2N + 1I$ incremental cycles, where S, N and I are as defined in **6.2 Cycle Types** on page 6-2.

STR instructions take 2N incremental cycles to execute.

ARM Instruction Set - LDR, STR

4.9.8 Assembler syntax

$$\langle \text{LDR} \mid \text{STR} \rangle \{ \text{cond} \} \{ \text{B} \} \{ \text{T} \} \text{ Rd}, \langle \text{Address} \rangle$$

where:

LDR load from memory into a register

STR store from a register into memory

{cond} two-character condition mnemonic. See ► *Table 4-2: Condition code summary* on page 4-5.

{B} if B is present then byte transfer, otherwise word transfer

{T} if T is present the W bit will be set in a post-indexed instruction, forcing non-privileged mode for the transfer cycle. T is not allowed when a pre-indexed addressing mode is specified or implied.

Rd is an expression evaluating to a valid register number.

Rn and Rm are expressions evaluating to a register number. If Rn is R15 then the assembler will subtract 8 from the offset value to allow for ARM7TDMI pipelining. In this case base write-back should not be specified.

<Address> can be:

- 1 An expression which generates an address:

<expression>

The assembler will attempt to generate an instruction using the PC as a base and a corrected immediate offset to address the location given by evaluating the expression. This will be a PC relative, pre-indexed address. If the address is out of range, an error will be generated.

- 2 A pre-indexed addressing specification:

| [Rn] | offset of zero |
|--------|----------------|
| 0 | 0 |
| 1 | 1 |
| 2 | 2 |
| 3 | 3 |
| 4 | 4 |
| 5 | 5 |
| 6 | 6 |
| 7 | 7 |
| 8 | 8 |
| 9 | 9 |
| 10 | 10 |
| 11 | 11 |
| 12 | 12 |
| 13 | 13 |
| 14 | 14 |
| 15 | 15 |
| 16 | 16 |
| 17 | 17 |
| 18 | 18 |
| 19 | 19 |
| 20 | 20 |
| 21 | 21 |
| 22 | 22 |
| 23 | 23 |
| 24 | 24 |
| 25 | 25 |
| 26 | 26 |
| 27 | 27 |
| 28 | 28 |
| 29 | 29 |
| 30 | 30 |
| 31 | 31 |
| 32 | 32 |
| 33 | 33 |
| 34 | 34 |
| 35 | 35 |
| 36 | 36 |
| 37 | 37 |
| 38 | 38 |
| 39 | 39 |
| 40 | 40 |
| 41 | 41 |
| 42 | 42 |
| 43 | 43 |
| 44 | 44 |
| 45 | 45 |
| 46 | 46 |
| 47 | 47 |
| 48 | 48 |
| 49 | 49 |
| 50 | 50 |
| 51 | 51 |
| 52 | 52 |
| 53 | 53 |
| 54 | 54 |
| 55 | 55 |
| 56 | 56 |
| 57 | 57 |
| 58 | 58 |
| 59 | 59 |
| 60 | 60 |
| 61 | 61 |
| 62 | 62 |
| 63 | 63 |
| 64 | 64 |
| 65 | 65 |
| 66 | 66 |
| 67 | 67 |
| 68 | 68 |
| 69 | 69 |
| 70 | 70 |
| 71 | 71 |
| 72 | 72 |
| 73 | 73 |
| 74 | 74 |
| 75 | 75 |
| 76 | 76 |
| 77 | 77 |
| 78 | 78 |
| 79 | 79 |
| 80 | 80 |
| 81 | 81 |
| 82 | 82 |
| 83 | 83 |
| 84 | 84 |
| 85 | 85 |
| 86 | 86 |
| 87 | 87 |
| 88 | 88 |
| 89 | 89 |
| 90 | 90 |
| 91 | 91 |
| 92 | 92 |
| 93 | 93 |
| 94 | 94 |
| 95 | 95 |
| 96 | 96 |
| 97 | 97 |
| 98 | 98 |
| 99 | 99 |
| 100 | 100 |
| 101 | 101 |
| 102 | 102 |
| 103 | 103 |
| 104 | 104 |
| 105 | 105 |
| 106 | 106 |
| 107 | 107 |
| 108 | 108 |
| 109 | 109 |
| 110 | 110 |
| 111 | 111 |
| 112 | 112 |
| 113 | 113 |
| 114 | 114 |
| 115 | 115 |
| 116 | 116 |
| 117 | 117 |
| 118 | 118 |
| 119 | 119 |
| 120 | 120 |
| 121 | 121 |
| 122 | 122 |
| 123 | 123 |
| 124 | 124 |
| 125 | 125 |
| 126 | 126 |
| 127 | 127 |
| 128 | 128 |
| 129 | 129 |
| 130 | 130 |
| 131 | 131 |
| 132 | 132 |
| 133 | 133 |
| 134 | 134 |
| 135 | 135 |
| 136 | 136 |
| 137 | 137 |
| 138 | 138 |
| 139 | 139 |
| 140 | 140 |
| 141 | 141 |
| 142 | 142 |
| 143 | 143 |
| 144 | 144 |
| 145 | 145 |
| 146 | 146 |
| 147 | 147 |
| 148 | 148 |
| 149 | 149 |
| 150 | 150 |
| 151 | 151 |
| 152 | 152 |
| 153 | 153 |
| 154 | 154 |
| 155 | 155 |
| 156 | 156 |
| 157 | 157 |
| 158 | 158 |
| 159 | 159 |
| 160 | 160 |
| 161 | 161 |
| 162 | 162 |
| 163 | 163 |
| 164 | 164 |
| 165 | 165 |
| 166 | 166 |
| 167 | 167 |
| 168 | 168 |
| 169 | 169 |
| 170 | 170 |

| | |
|-----------------------|---------------------------------|
| [Rn,<#expression>]{!} | offset of <expression> bytes |
|-----------------------|---------------------------------|

| | |
|-----------------------------------|--|
| [Rn, {+/-}Rm{ , <shift> }] { ! } | offset of +/- contents of index register, shifted by <shift> |
|-----------------------------------|--|

- 3 A post-indexed addressing specification:

| | |
|---------------------|---------------------------------|
| [Rn], <#expression> | offset of <expression> bytes |
|---------------------|---------------------------------|

| | |
|-------------------------|--|
| [Rn], {+/-}Rm{,<shift>} | offset of +/- contents of index register, shifted as by <shift>. |
|-------------------------|--|

<shift> general shift operation (see data processing instructions) but you cannot specify the shift amount by a register.
 {!} writes back the base register (set the W bit) if! is present.

4.9.9 Examples

```
STR    R1,[R2,R4]!           ; Store R1 at R2+R4 (both of which are
                              ; registers) and write back address to
                              ; R2.
STR    R1,[R2],R4            ; Store R1 at R2 and write back
                              ; R2+R4 to R2.
LDR    R1,[R2,#16]           ; Load R1 from contents of R2+16, but
                              ; don't write back.
LDR    R1,[R2,R3,LSL#2]      ; Load R1 from contents of R2+R3*4.
LDREQBR1,[R6,#5]            ; Conditionally load byte at R6+5 into
                              ; R1 bits 0 to 7, filling bits 8 to 31
                              ; with zeros.
STR    R1,PLACE              ; Generate PC relative offset to
                              ; address PLACE.
•
PLACE
```

ARM Instruction Set - LDM, STM

4.11 Block Data Transfer (LDM, STM)

The instruction is only executed if the condition is true. The various conditions are defined in [Table 4-2: Condition code summary](#) on page 4-5. The instruction encoding is shown in [Figure 4-18: Block data transfer instructions](#).

Block data transfer instructions are used to load (LDM) or store (STM) any subset of the currently visible registers. They support all possible stacking modes, maintaining full or empty stacks which can grow up or down memory, and are very efficient instructions for saving or restoring context, or for moving large blocks of data around main memory.

4.11.1 The register list

The instruction can cause the transfer of any registers in the current bank (and non-user mode programs can also transfer to and from the user bank, see below). The register list is a 16 bit field in the instruction, with each bit corresponding to a register. A 1 in bit 0 of the register field will cause R0 to be transferred, a 0 will cause it not to be transferred; similarly bit 1 controls the transfer of R1, and so on.

Any subset of the registers, or all the registers, may be specified. The only restriction is that the register list should not be empty.

Whenever R15 is stored to memory the stored value is the address of the STM instruction plus 12.




Figure 4-18: Block data transfer instructions

4.11.2 Addressing modes

The transfer addresses are determined by the contents of the base register (Rn), the pre/post bit (P) and the up/down bit (U). The registers are transferred in the order lowest to highest, so R15 (if in the list) will always be transferred last. The lowest register also gets transferred to/from the lowest memory address. By way of illustration, consider the transfer of R1, R5 and R7 in the case where Rn=0x1000 and write back of the modified base is required (W=1). *Figure 4-19: Post-increment addressing*, *Figure 4-20: Pre-increment addressing*, *Figure 4-21: Post-decrement addressing* and *Figure 4-22: Pre-decrement addressing* show the sequence of register transfers, the addresses used, and the value of Rn after the instruction has completed.

In all cases, had write back of the modified base not been required (W=0), Rn would have retained its initial value of 0x1000 unless it was also in the transfer list of a load multiple register instruction, when it would have been overwritten with the loaded value.

4.11.3 Address alignment

The address should normally be a word aligned quantity and non-word aligned addresses do not affect the instruction. However, the bottom 2 bits of the address will appear on **A[1:0]** and might be interpreted by the memory system.




Figure 4-19: Post-increment addressing

ARM Instruction Set - LDM, STM




Figure 4-20: Pre-increment addressing




Figure 4-21: Post-decrement addressing




Figure 4-22: Pre-decrement addressing

4.11.4 Use of the S bit

When the S bit is set in a LDM/STM instruction its meaning depends on whether or not R15 is in the transfer list and on the type of instruction. The S bit should only be set if the instruction is to execute in a privileged mode.

LDM with R15 in transfer list and S bit set (Mode changes)

If the instruction is a LDM then SPSR_<mode> is transferred to CPSR at the same time as R15 is loaded.

STM with R15 in transfer list and S bit set (User bank transfer)

The registers transferred are taken from the User bank rather than the bank corresponding to the current mode. This is useful for saving the user state on process switches. Base write-back should not be used when this mechanism is employed.

R15 not in list and S bit set (User bank transfer)

For both LDM and STM instructions, the User bank registers are transferred rather than the register bank corresponding to the current mode. This is useful for saving the user state on process switches. Base write-back should not be used when this mechanism is employed.

When the instruction is LDM, care must be taken not to read from a banked register during the following cycle (inserting a dummy instruction such as MOV R0, R0 after the LDM will ensure safety).

ARM Instruction Set - LDM, STM

4.11.5 Use of R15 as the base

R15 should not be used as the base register in any LDM or STM instruction.

4.11.6 Inclusion of the base in the register list

When write-back is specified, the base is written back at the end of the second cycle of the instruction. During a STM, the first register is written out at the start of the second cycle. A STM which includes storing the base, with the base as the first register to be stored, will therefore store the unchanged value, whereas with the base second or later in the transfer order, will store the modified value. A LDM will always overwrite the updated base if the base is in the list.

4.11.7 Data aborts

Some legal addresses may be unacceptable to a memory management system, and the memory manager can indicate a problem with an address by taking the **ABORT** signal HIGH. This can happen on any transfer during a multiple register load or store, and must be recoverable if ARM7TDMI is to be used in a virtual memory system.

Aborts during STM instructions

If the abort occurs during a store multiple instruction, ARM7TDMI takes little action until the instruction completes, whereupon it enters the data abort trap. The memory manager is responsible for preventing erroneous writes to the memory. The only change to the internal state of the processor will be the modification of the base register if write-back was specified, and this must be reversed by software (and the cause of the abort resolved) before the instruction may be retried.

Aborts during LDM instructions

When ARM7TDMI detects a data abort during a load multiple instruction, it modifies the operation of the instruction to ensure that recovery is possible.

- 1 Overwriting of registers stops when the abort happens. The aborting load will not take place but earlier ones may have overwritten registers. The PC is always the last register to be written and so will always be preserved.
- 2 The base register is restored, to its modified value if write-back was requested. This ensures recoverability in the case where the base register is also in the transfer list, and may have been overwritten before the abort occurred.

The data abort trap is taken when the load multiple has completed, and the system software must undo any base modification (and resolve the cause of the abort) before restarting the instruction.

4.11.8 Instruction cycle times

Normal LDM instructions take $nS + 1N + 1I$ and LDM PC takes $(n+1)S + 2N + 1I$ incremental cycles, where S, N and I are as defined in **6.2 Cycle Types** on page 6-2. STM instructions take $(n-1)S + 2N$ incremental cycles to execute, where n is the number of words transferred.

4.11.9 Assembler syntax

`<LDM|STM> {cond} <FD|ED|FA|EA|IA|IB|DA|DB> Rn{!}, <Rlist>{^}`

where:

- `{cond}` two character condition mnemonic. See [Table 4-2: Condition code summary](#) on page 4-5.
- `Rn` is an expression evaluating to a valid register number
- `<Rlist>` is a list of registers and register ranges enclosed in {} (e.g. {R0,R2-R7,R10}).
- `{!}` if present requests write-back (W=1), otherwise W=0
- `{^}` if present set S bit to load the CPSR along with the PC, or force transfer of user bank when in privileged mode

Addressing mode names

There are different assembler mnemonics for each of the addressing modes, depending on whether the instruction is being used to support stacks or for other purposes. The equivalence between the names and the values of the bits in the instruction are shown in the following table:

| Name | Stack | Other | L bit | P bit | U bit |
|----------------------|-------|-------|-------|-------|-------|
| pre-increment load | LDMED | LDMIB | 1 | 1 | 1 |
| post-increment load | LDMFD | LDMIA | 1 | 0 | 1 |
| pre-decrement load | LDMEA | LDMDB | 1 | 1 | 0 |
| post-decrement load | LDMFA | LDMDA | 1 | 0 | 0 |
| pre-increment store | STMFA | STMIB | 0 | 1 | 1 |
| post-increment store | STMEA | STMIA | 0 | 0 | 1 |
| pre-decrement store | STMFD | STMDB | 0 | 1 | 0 |
| post-decrement store | STMED | STMDA | 0 | 0 | 0 |

Table 4-6: Addressing mode names

FD, ED, FA, EA define pre/post indexing and the up/down bit by reference to the form of stack required. The F and E refer to a “full” or “empty” stack, i.e. whether a pre-index has to be done (full) before storing to the stack. The A and D refer to whether the stack is ascending or descending. If ascending, a STM will go up and LDM down, if descending, vice-versa.

IA, IB, DA, DB allow control when LDM/STM are not being used for stacks and simply mean Increment After, Increment Before, Decrement After, Decrement Before.

ARM Instruction Set - LDM, STM

4.11.10 Examples

```
LDMFD SP!, {R0,R1,R2}      ; Unstack 3 registers.
STMIA R0, {R0-R15}         ; Save all registers.
LDMFD SP!, {R15}           ; R15 <- (SP), CPSR unchanged.
LDMFD SP!, {R15}^          ; R15 <- (SP), CPSR <- SPSR_mode
                           ; (allowed only in privileged modes).
STMFD R13, {R0-R14}^       ; Save user mode regs on stack
                           ; (allowed only in privileged modes).
```

These instructions may be used to save state on subroutine entry, and restore it efficiently on return to the calling routine:

```
STMED SP!, {R0-R3,R14}     ; Save R0 to R3 to use as workspace
                           ; and R14 for returning.
BL      somewhere          ; This nested call will overwrite R14
LDMED SP!, {R0-R3,R15}     ; restore workspace and return.
```

4.13 Software Interrupt (SWI)

The instruction is only executed if the condition is true. The various conditions are defined in [Table 4-2: Condition code summary](#) on page 4-5. The instruction encoding is shown in [Figure 4-24: Software interrupt instruction](#), below.




Figure 4-24: Software interrupt instruction

The software interrupt instruction is used to enter Supervisor mode in a controlled manner. The instruction causes the software interrupt trap to be taken, which effects the mode change. The PC is then forced to a fixed value (0x08) and the CPSR is saved in SPSR_svc. If the SWI vector address is suitably protected (by external memory management hardware) from modification by the user, a fully protected operating system may be constructed.

4.13.1 Return from the supervisor

The PC is saved in R14_svc upon entering the software interrupt trap, with the PC adjusted to point to the word after the SWI instruction. MOVS PC,R14_svc will return to the calling program and restore the CPSR.

Note that the link mechanism is not re-entrant, so if the supervisor code wishes to use software interrupts within itself it must first save a copy of the return address and SPSR.

4.13.2 Comment field

The bottom 24 bits of the instruction are ignored by the processor, and may be used to communicate information to the supervisor code. For instance, the supervisor may look at this field and use it to index into an array of entry points for routines which perform the various supervisor functions.

4.13.3 Instruction cycle times

Software interrupt instructions take $2S + 1N$ incremental cycles to execute, where S and N are as defined in [6.2 Cycle Types](#) on page 6-2.

ARM Instruction Set - SWI

4.13.4 Assembler syntax

| | |
|------------------------|--|
| SWI{cond} <expression> | |
| {cond} | two character condition mnemonic, Table 4-2: Condition code summary on page 4-5. |
| <expression> | is evaluated and placed in the comment field (which is ignored by ARM7TDMI). |

4.13.5 Examples

```
SWI    ReadC                ; Get next character from read stream.
SWI    WriteI+"k"           ; Output a "k" to the write stream.
SWINE  0                    ; Conditionally call supervisor
                          ; with 0 in comment field.
```

Supervisor code

The previous examples assume that suitable supervisor code exists, for instance:

```
0x08 B Supervisor          ; SWI entry point
EntryTable                  ; addresses of supervisor routines
    DCD ZeroRtn
    DCD ReadCRtn
    DCD WriteIRtn
    . . .
Zero    EQU    0
ReadC   EQU    256
WriteI  EQU    512

Supervisor

; SWI has routine required in bits 8-23 and data (if any) in
; bits 0-7.
; Assumes R13_svc points to a suitable stack

STMFD R13,{R0-R2,R14}      ; Save work registers and return
                          ; address.
LDR    R0,[R14,#-4]         ; Get SWI instruction.
BIC    R0,R0,#0xFF000000    ; Clear top 8 bits.
MOV    R1,R0,LSR#8          ; Get routine offset.
ADR    R2,EntryTable        ; Get start address of entry table.
LDR    R15,[R2,R1,LSL#2]    ; Branch to appropriate routine.

    WriteIRtn               ; Enter with character in R0 bits 0-7.
    . . .
LDMFD R13,{R0-R2,R15}^     ; Restore workspace and return,
                          ; restoring processor mode and flags.
```


ARM Instruction Set - Undefined

4.17 Undefined Instruction

The instruction is only executed if the condition is true. The various conditions are defined in **Table 4-2: Condition code summary** on page 4-5. The instruction format is shown in **Figure 4-28: Undefined instruction**.




Figure 4-28: Undefined instruction

If the condition is true, the undefined instruction trap will be taken.

Note that the undefined instruction mechanism involves offering this instruction to any coprocessors which may be present, and all coprocessors must refuse to accept it by driving **CPA** and **CPB** HIGH.

4.17.1 Instruction cycle times

This instruction takes $2S + 1I + 1N$ cycles, where S, N and I are as defined in **►6.2 Cycle Types** on page 6-2.

4.17.2 Assembler syntax

The assembler has no mnemonics for generating this instruction. If it is adopted in the future for some specified use, suitable mnemonics will be added to the assembler. Until such time, this instruction must not be used.